

## Group No: 306 Bank Loan Status

First Name	Last Name	Email (hawk.iit.edu)	Student ID
Abhishek	Singh	asingh135@hawk.iit.edu	A20447125
Megha	Shrivastava	mshrivastava1@hawk.iit.edu	A20450886

## Table of Contents

1. Introduction.....	2
2. Data.....	2
3. System Flow of Project.....	4
4. Problem to be solved.....	4
5. Solutions.....	4
6. Experiment and Results.....	5
Data Cleaning.....	5
KNN.....	9
Logistic Model.....	12
Gradient Boosting.....	16
7. Conclusion.....	21

# 1. Introduction

Loan analysis is the method which helps in determining that the borrowers will be able to pay the loan or not. In this project we are taking 19 attributes, building a model and analyzing it further to check whether the person will be able to pay the loan or not. The dependent variable (Loan status) is a nominal variable we will convert that into numerical variable to apply linear regression model. This approach is also useful in understanding many things like which attribute will affect the loan delinquency. Also, the model analysis generally helps in determining and predicting the things which are related to measuring. The dataset contains the csv file credit\_train.csv which contains the various independent and dependent variables. We will use this information to analyze whether the person who takes the loan is able to pay it back or not with the help various attributes present in the dataset. The dataset contains credit score, annual income and years of job which is helpful in predicting whether the person can pay the loan or not. We have applied KNN and logistic regression model. Also, we have formed a gradient boosting model found the accuracy and also stated which model is best.

# 2. Data

The data set was found on the Kaggle which has 99981 records in the data. The dataset has two tables credit\_train.csv and credit\_test.csv. The data contains 19 columns and is present in the credit\_train.csv file. The data can be found from the link: <https://www.kaggle.com/zaurbegiev/my-dataset>

In the dataset except the loan status all are independent variable and the loan status is the dependent variable. By this data the analysis in between the dependent and independent variable is done and the further the graph is plot. There are only two values for the dependent variable which are “Fully paid” and “Charged off”.

### **Data understanding:**

The attributes present in the dataset are further divided into two categories: Qualitative and Quantitative.

Qualitative	Quantitative
Loan ID	Current Loan Amount
Customer ID	Credit Score
Loan Status	Annual Income
Term	Years in Current Job
Home Ownership	Monthly Debt
Purpose	Year of Credit History
	Month since last delinquent
	Number of Open Account
	Number of Credit Problems
	Current Credit Balance
	Maximum Open Credit
	Bankruptcies
	Tax Liens

Further the attributes are classified as:

- Qualitative:

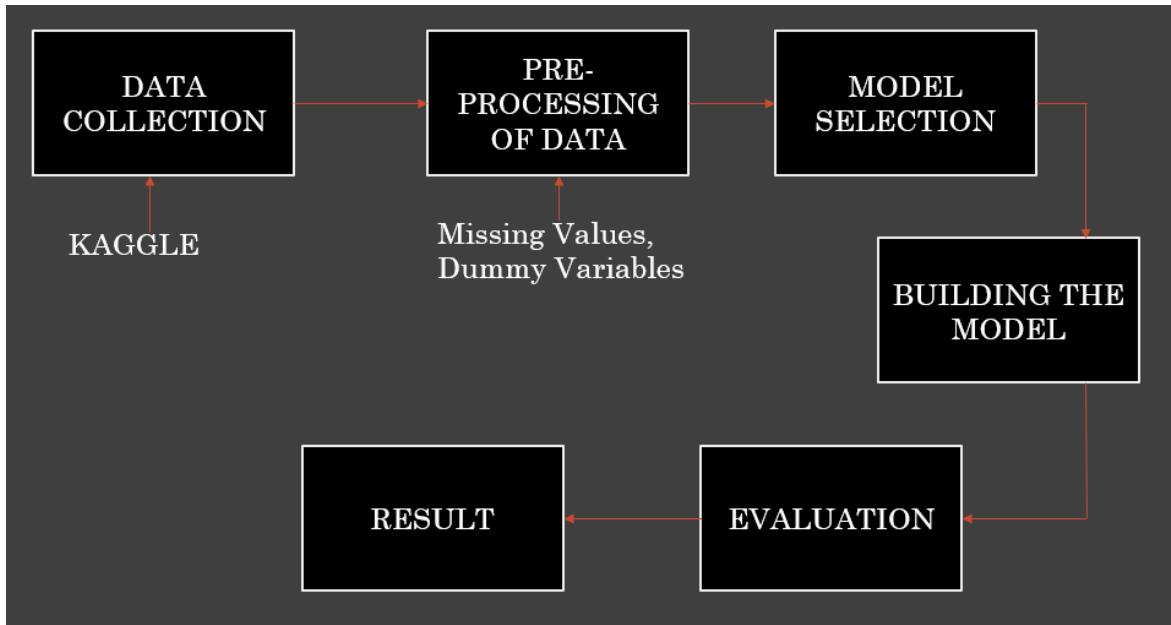
Loan ID: Nominal  
Customer ID: Nominal  
Loan Status: Nominal  
Term: Ordinal  
Home Ownership: Nominal  
Purpose: Nominal

➤ **Quantitative:**

Current Loan Amount: Discrete  
Credit Score: Discrete  
Annual Income: Discrete  
Years in Current Job: Discrete  
Monthly Debt: Discrete  
Year of Credit History: Discrete  
Month since last delinquent: Discrete  
Number of Open Account: Discrete  
Number of Credit Problems: Discrete  
Current Credit Balance: Discrete  
Maximum Open Credit: Discrete  
Bankruptcies: Discrete  
Tax Liens: Discrete

As we see according to the data that loan status is nominal variable, so we will change the data to binary form or the required form according to the analysis and then further analysis it. Like for logistic analysis we will make it to binary (0 or 1), for KNN we will change it to Yes or No. Also, all the remaining will also be converted according to the analysis.

### 3. System Flow of Project



### 4. Problem to be solved

By this we get know about the attributes which will affect whether any person will pay the loan or not. Also, we will build different model which will help in analyzing and predicting that the person will pay the loan or not according to the attributes. With the help of independent and dependent attribute present in the dataset, the problem that who will pay the loan is solved to a level. The data include a credit score column which shows and is helpful in providing the loan and also in predicting further outcomes

### 5. Solutions

#### Modeling:

- After preprocessing and normalizing the data we will build different models based on Hold-Out evaluation since our dataset has more than 5000 records.
- We will split the dataset into training and test data set.

We will create various models such as

A. KNN

B. Logistic Regression Model

C. Gradient Boosting Model

## 6. Experiments and Results

## Data Cleaning

In the below data set, you can see there are two columns Loan Id & Customer Id which is just for identification purpose so we will remove those columns. You can see in the next screen shot that we removed those columns.

**Loan Id & Customer Id in the Data set:**

	Loan.ID	Customer.ID	Loan.Status	Current.Loan.Amount	Term	Credit.Score
1	14dd8831-6af5-400b-83ec-68e61888a048	981165ec-3274-42f5-a3b4-d104041a9ca9	Fully Paid		445412	Short Term
2	4771cc26-131a-45db-b5aa-537ea4ba5342	2de017a3-2e01-49cb-a581-08169e83be29	Fully Paid		262328	Short Term
3	4eed4e6a-aa2f-4c91-8651-ce984ee8fb26	5efb2b2b-bf11-4dfd-a572-3761a2694725	Fully Paid		99999999	Short Term
4	77598f7b-32e7-4e3b-a6e5-06ba0d98fe8a	e777faab-98ae-45af-9a86-7ce5b33b1011	Fully Paid		347666	Long Term
5	d4062e70-befa-4995-8643-a0de73938182	81536ad9-5ccf-4eb8-befb-47a4d608658e	Fully Paid		176220	Short Term
6	89d8cb0c-e5c2-4f54-b056-48a645c543dd	4ffe99d3-7f2a-44db-afc1-40943f1f9750	Charged Off		206602	Short Term
	Annual.Income	Years.in.current.job	Home.Ownership	Purpose	Monthly.Debt	Years.of.Credit.History
1	1167493	8 years	Home Mortgage	Home Improvements	5214.74	17.2
2	NA	10+ years	Home Mortgage	Debt Consolidation	33295.98	21.1
3	2231892	8 years	Own Home	Debt Consolidation	29200.53	14.9
4	806949	3 years	Own Home	Debt Consolidation	8741.90	12.0
5	NA	5 years	Rent	Debt Consolidation	20639.70	6.1
6	896857	10+ years	Home Mortgage	Debt Consolidation	16367.74	17.3
	Months.since.last.delinquent	Number.of.Open.Accounts	Number.of.Credit.Problems	Current.Credit.Balance	Maximum.Open.Credit	
1	NA	6	1	228190	416746	
2	8	35	0	229976	850784	
3	29	18	1	297996	750090	
4	NA	9	0	256329	386958	
5	NA	15	0	253460	427174	
6	NA	6	0	215308	272448	
	Bankruptcies	Tax.Liens				
1	1	0				
2	0	0				
3	0	0				
4	0	0				
5	0	0				
6	0	0				

**Removing the Loan id & Customer Id column from the Data set:**

We have created a function to calculate the null values in the data set for all the columns, you can see the output below. As you can see there are 12 columns which consist N/A values. We can also see the number of N/A values too.

### **Function to calculate null values:**

```
> # code for finding N/A values in the data set.  
> missing_data <- colSums(is.na(credit))[colSums(is.na(credit)) > 0] %>% sort(decreasing=T)  
> missing_data
```

Months.since.last.delinquent	Credit.Score	Annual.Income
53655	19668	19668
Bankruptcies	Tax.Liens	Maximum.Open.Credit
718	524	516
Current.Loan.Amount	Monthly.Debt	Years.of.Credit.History
514	514	514
Number.of.Open.Accounts	Number.of.Credit.Problems	current.Credit.Balance
514	514	514

Since the Months.since.last delinquent column contains more than 50% of null values. We have removed that column in below screen shot:

```

> #removing Months.since.last.delinquent column as it contains more than 50% of N/A value
> credit=select(credit,-"Months.since.last.delinquent")
> head(credit)
   Loan.Status Current.Loan.Amount      Term Credit.Score Annual.Income Years.in.current.job Home.Ownership
1 Fully Paid           445412 Short Term       709     1167493          8 years Home Mortgage
2 Fully Paid           262328 Short Term        NA        NA          10+ years Home Mortgage
3 Fully Paid         999999999 Short Term       741     2231892          8 years Own Home
4 Fully Paid           347666 Long Term        721     806949          3 years Own Home
5 Fully Paid           176220 Short Term        NA        NA          5 years       Rent
6 Charged off          206602 Short Term       7290     896857          10+ years Home Mortgage
   Purpose Monthly.Debt Years.of.Credit.History Number.of.Open.Accounts Number.of.Credit.Problems
1 Home Improvements      5214.74            17.2                  6                      1
2 Debt Consolidation    33295.98            21.1                 35                      0
3 Debt Consolidation    29200.53            14.9                 18                      1
4 Debt Consolidation    8741.90             12.0                  9                      0
5 Debt Consolidation   20639.70              6.1                 15                      0
6 Debt Consolidation   16367.74            17.3                  6                      0
   Current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
1                   228190            416746             1                  0
2                   229976            850784             0                  0
3                   297996            750090             0                  0
4                   256329            386958             0                  0
5                   253460            427174             0                  0
6                   215308            272448             0                  0

```

You can see in the below screen shot there are 514 rows at the tail of the data set that contains all the N/A values. So, we have removed all those rows.

We can see now at the bottom of the data set that there no N/A Rows:

```

> credit<- credit[-seq(nrow(credit),nrow(credit)-514),]
> tail(credit)
   Loan.Status Current.Loan.Amount Term Credit.Score Annual.Income Years.in.current.job Home.Ownership
99994 Fully Paid          44484 Short Term      717     1152426           10+ years    Home Mortgage
99995 Fully Paid          210584 short Term      719      783389            1 year     Home Mortgage
99996 Fully Paid          147070 Short Term      725      475437             7 years    Own Home
99997 Fully Paid          99999999 short Term      732     1289416            1 year     Rent
99998 Fully Paid          103136 Short Term      742     1150545            6 years     Rent
99999 Fully Paid          530332 Short Term      746     1717524            9 years     Rent
                                         Purpose Monthly.Debt Years.of.Credit.History Number.of.Open.Accounts Number.of.Credit.Problems
99994 small_business        6280.64           21.0                   6                         0
99995 other                  3727.61           17.4                   6                         0
99996 other                  2202.86           22.3                   5                         0
99997 Debt Consolidation  13109.05           9.4                  22                         0
99998 Debt Consolidation  7315.57            18.8                  12                         1
99999 Debt Consolidation  9890.07            15.0                   8                         0
   Current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
99994                      961932                 0            0            0
99995                        456            259160            0            0
99996                      477666            658548            0            0
99997                      153045            509234            0            0
99998                      109554            537548            1            0
99999                      404225            738254            0            0
> |

```

**Now, we are left with the 5 columns which contains N/A values. We can see that in the below screen shot:**

```

> missing_data <- colsums(is.na(credit))[colsums(is.na(credit)) > 0] %>% sort(decreasing=T)
> missing_data
  Credit.Score      Annual.Income      Bankruptcies      Tax.Liens Maximum.Open.Credit
1       19154          19154             204                 10                  2
  |

```

To remove those N/A values from the 5 columns. We replaced those values with the mean value of the dataset. The code for that is given below.

```

#swaping the N/A values with the mean value
credit$Credit.Score=ifelse(is.na(credit$Credit.Score),ave(credit$Credit.Score, FUN= function(x) mean(x,na.rm=TRUE)),credit$Credit.Score)
credit$Annual.Income=ifelse(is.na(credit$Annual.Income),ave(credit$Annual.Income, FUN= function(x) mean(x,na.rm=TRUE)),credit$Annual.Income)
credit$Bankruptcies=ifelse(is.na(credit$Bankruptcies),ave(credit$Bankruptcies, FUN= function(x) mean(x,na.rm=TRUE)),credit$Bankruptcies)
credit$Tax.Liens=ifelse(is.na(credit$Tax.Liens),ave(credit$Tax.Liens, FUN= function(x) mean(x,na.rm=TRUE)),credit$Tax.Liens)
credit$Maximum.Open.Credit=ifelse(is.na(credit$Maximum.Open.Credit),ave(credit$Maximum.Open.Credit, FUN= function(x) mean(x,na.rm=TRUE)),credit$Maximum.Open.Credit)
  |

```

### No N/A values present:

```

> missing_data <- colsums(is.na(credit))[colsums(is.na(credit)) > 0] %>% sort(decreasing=T)
> missing_data
named numeric(0)
  |

```

### Dataset with No N/A values:

```

> head(credit,n=10)
  Loan.Status Current.Loan.Amount      Term Credit.Score Annual.Income Years.in.current.job Home.Ownership Purpose Monthly.Debt
1 Fully Paid           445412 Short Term    709.00     1167493            8 years   Home Mortgage   Home Improvements    5214.74
2 Fully Paid           262328 Short Term   1076.46     1378282           10+ years  Home Mortgage Debt Consolidation 33295.98
3 Fully Paid          99999999 Short Term    741.00     2231892            8 years   Own Home Debt Consolidation 29200.53
4 Fully Paid           347666 Long Term   721.00      806949            3 years   Own Home Debt Consolidation 8741.90
5 Fully Paid           176220 Short Term   1076.46     1378282            5 years   Rent Debt Consolidation 20639.70
6 Charged off         206602 Short Term   7290.00     896857           10+ years  Home Mortgage Debt Consolidation 16367.74
7 Fully Paid           217646 Short Term   730.00     1184194            < 1 year  Home Mortgage Debt Consolidation 10855.08
8 Charged off          648714 Long Term   1076.46     1378282            < 1 year  Home Mortgage        Buy House 14806.13
9 Fully Paid           548746 Short Term   678.00      2559110            2 years   Rent Debt Consolidation 18660.28
10 Fully Paid          215952 Short Term   739.00     1454735            < 1 year  Rent Debt Consolidation 39277.75
  Years.of.Credit.History Number.of.Open.Accounts Number.of.Credit.Problems current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
1                   17.2                      6                     1           228190        416746           1          0
2                   21.1                      35                    0           229976        850784           0          0
3                   14.9                      18                    1           297996        750090           0          0
4                   12.0                      9                     0           256329        386958           0          0
5                   6.1                      15                    0           253460        427174           0          0
6                   17.3                      6                     0           215308        272448           0          0
7                   19.6                      13                    1           122170        272052           1          0
8                   8.2                      15                    0           193306        864204           0          0
  |

```

## A. KNN

- KNN is a supervised learning model. We must have predefined labels for this model. In the given data set we have predefined labels for all the attributes.
- We also should have knowledge from the past data for this model.

- There are three types of classification in this model. We will use binary classification because our dependent variable (loan status) has only two values which are “Fully paid” or “charged off”.

KNN classifier:

- In this algorithm we calculate the distance between the target (Loan Status) and various instances.
- Then we will identify the k-nearest neighbor.
- Then we will predict the labels and validate with the truth. We can use Manhattan distance or Euclidian distance to calculate the distance between the target (Loan Status) and various instances

## Normalization

To calculate the distance between the dependent variable and each independent variable. We need to first normalize all the numeric values. So that we can correctly calculate the distance. We used min-max function to calculate the relative distance the minimum and maximum value of the column.

```
#normalize selected data using function scale
credit[num.vars] <- lapply(credit[num.vars], scale)
credit[num.vars]

head(credit)
credit[num.vars] <- apply(credit[num.vars], 2, FUN = function(x) (x - min(x))/(max(x)-min(x)))
head(credit[num.vars])

head(credit)
```

## Normalized data set:

```
head(credit)
# Loan.Status Current.Loan.Amount Term Credit.Score Annual.Income Years.in.current.job Home.Ownership Purpose Monthly.Debt
# Fully Paid 0.004346570 Short Term 0.01790614 0.006592101 8 years Home Mortgage Home Improvements 0.01196471
# Fully Paid 0.002515532 Short Term 0.07096898 0.007865899 10+ years Home Mortgage Debt Consolidation 0.07639439
# Fully Paid 1.000000000 Short Term 0.0252708 0.013024263 8 years Own Home Debt Consolidation 0.06699777
# Fully Paid 0.003369004 Long Term 0.01963899 0.004413335 3 years Own Home Debt Consolidation 0.02005744
# Fully Paid 0.001654359 Short Term 0.07096898 0.007865899 5 years Rent Debt Consolidation 0.04735578
# Charged off 0.001958212 Short Term 0.96823105 0.004956649 10+ years Home Mortgage Debt Consolidation 0.03755419
# Years.of.Credit.History Number.of.Open.Accounts Number.of.Credit.Problems Current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
# 0.20328849 0.07894737 0.06666667 0.006940303 0.0002706604 0.1428571 0
# 0.26158445 0.46052632 0.00000000 0.006994623 0.0005525512 0.0000000 0
# 0.16890882 0.23684211 0.06666667 0.009063423 0.0004871543 0.0000000 0
# 0.12556054 0.11842105 0.00000000 0.007796139 0.0002513142 0.0000000 0
# 0.03736921 0.19736842 0.00000000 0.007708879 0.0002774329 0.0000000 0
# 0.20478326 0.07894737 0.00000000 0.006548502 0.0001769444 0.0000000 0
```

## Creating Dummies:

As we know we can normalize the numeric columns using the min-max function. Although, we cannot do this for categorical columns. So, we create dummy variables for categorical data values.

```
#dummies
install.packages("dummies")
library(dummies)

credit.dummy=dummy.data.frame(credit,names=c("Term","Years.in.current.job","Home.Ownership","Purpose"))

head(credit.dummy)

#Removing the extra column
```

## Data set with dummy variables:

```
head(credit.dummy)
Loan.Status Current.Loan.Amount TermLong Term TermShort Term Credit.Score Annual.Income Years.in.current.job< 1 year Years.in.current.job1 year
Fully Paid 0.004346570 0 1 0.01790614 0.006592101 0 0
Fully Paid 0.002515532 0 1 0.07096898 0.007865899 0 0
Fully Paid 1.000000000 0 1 0.02252708 0.013024263 0 0
Fully Paid 0.003369004 1 0 0.01963899 0.004413335 0 0
Fully Paid 0.001654359 0 1 0.07096898 0.007865899 0 0
Charged off 0.001958212 0 1 0.96823105 0.004956649 0 0
Years.in.current.job10+ years Years.in.current.job2 years Years.in.current.job3 years Years.in.current.job4 years Years.in.current.job5 years
0 0 0 0 0 0 0
1 0 0 0 0 0 0
0 0 0 0 0 0 0
0 0 0 1 0 0 0
0 0 0 0 0 0 0
1 0 0 0 0 0 1
Years.in.current.job6 years Years.in.current.job7 years Years.in.current.job8 years Years.in.current.job9 years Years.in.current.jobn/a
0 0 0 1 0 0
0 0 0 0 0 0
0 0 0 1 0 0
0 0 0 0 0 0
0 0 0 0 0 0
0 0 0 0 0 0
```

## Removing the extra column:

We only need n-1 levels for each column. So, we will remove the extra dummy variable column from the data set.

```
#Removing the extra column
credit.final=select(credit.dummy,-c("TermLong Term","Years.in.current.jobn/a","Home.OwnershipHaveMortgage","PurposeBusiness Loan"))
head(credit.final)

Loan.Status Current.Loan.Amount TermShort Term Credit.Score Annual.Income Years.in.current.job< 1 year Years.in.current.job1 year Years.in.current.job10+ years
1 Fully Paid 0.004346570 1 0.01790614 0.006592101 0 0 0
2 Fully Paid 0.002515532 1 0.07096898 0.007865899 0 0 1
3 Fully Paid 1.000000000 1 0.02252708 0.013024263 0 0 0
4 Fully Paid 0.003369004 0 0.01963899 0.004413335 0 0 0
5 Fully Paid 0.001654359 1 0.07096898 0.007865899 0 0 0
6 Charged off 0.001958212 1 0.96823105 0.004956649 0 0 1
Years.in.current.job2 years Years.in.current.job3 years Years.in.current.job4 years Years.in.current.job5 years Years.in.current.job6 years
1 0 0 0 0 0 0
2 0 0 0 0 0 0
3 0 0 0 0 0 0
4 0 1 0 0 0 0
5 0 0 0 0 1 0
6 0 0 0 0 0 0
Years.in.current.job7 years Years.in.current.job8 years Years.in.current.job9 years Home.OwnershipHome Mortgage Home.OwnershipOwn Home.OwnershipRent
1 0 1 0 1 0 0
2 0 0 0 1 0 0
3 0 1 0 0 1 0
4 0 0 0 0 1 0
5 0 0 0 0 0 1
6 0 0 0 1 0 0
```

## KNN Model

We have created KNN model using N fold cross validation to determine the best K value for our data set. After creating the model, we can see that we got a best accuracy on K=24 which is 0.8119581= 81.19%

```

> model.nfold1 = train(credit.final[,2:42],credit.final$Loan.Status,'knn'
id(k = 20:30))
> model.nfold1
k-Nearest Neighbors

99999 samples
 41 predictor
 2 classes: '0', '1'

No pre-processing
Resampling: Cross-validated (10 fold)
Summary of sample sizes: 89999, 89999, 89999, 90000, 89999, 89999, ...
Resampling results across tuning parameters:

  k    Accuracy   Kappa
 20  0.8106481  0.2553137
 21  0.8113581  0.2553006
 22  0.8110681  0.2533087
 23  0.8113781  0.2526653
 24  0.8119581  0.2548797
 25  0.8116681  0.2513339
 26  0.8114181  0.2500168
 27  0.8115381  0.2489316
 28  0.8116081  0.2487669
 29  0.8115781  0.2469254
 30  0.8116581  0.2467509

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 24.
>

```

## Model for K=24:

```

library(class)
#final model KNN
# training and testing data
credit.final=credit.final[sample(nrow(credit.final)),]
select.data = sample (1:nrow(credit.final), 0.80*nrow(credit.final))
train.data.final=credit.final[select.data,]
test.data.final = credit.final[-select.data,]
head(train.data.final)
head(test.data.final)

#creating labels for training data and testing data
train.Loan.status.final=train.data.final$Loan.Status
head(train.Loan.status.final)
train.data.final=select(train.data.final,-1)
head(train.data.final)

test.Loan.status.final=test.data.final$Loan.Status
head(test.Loan.status.final)
test.data.final=select(test.data.final,-1)
head(test.data.final)

#KNN model
model.final=knn(train = train.data.final,test = test.data.final,cl=train.Loan.status.final,k=24)
summary(model.final)

```

## Confusion Metrics:

```
> table(model.final,test.Loan.status.final)
   test.Loan.status.final
model.final      Charged Off Fully Paid
                  0          0        0
Charged off     0        1468      1420
Fully Paid      0        2986      14126
> |
```

After creating confusion metrics. We will calculate Precision and Recall:

Precision: True positive/( True positive + False Positive)=1468/(1468+1420)

Precision = 0.5083=50.83%

Recall: True Positive/(True Positive + False Negative)= 1468/(1468+2986)

Recall= 0.3295=32.95%

## B.Logistic Regression Model

- Logistic Regression can be applied for binary dependent variable and various independent variables.
- In the dataset, dependent variable is loan Status which has value “Fully paid” or “Charged off”. So, we can apply the logistic regression model because we have only two values for it.
- We will use different feature selection methods such as backward, forward, stepwise etc. We will select the best model based on the AIC value.

We will perform all the data cleaning techniques which we performed the for the KNN. We can see the data below.

### Data Set:

```
> tail(credit.log)
   Loan.Status Current.Loan.Amount Term Credit.Score Annual.Income Years.in.current.job Home.Ownership
99994 Fully Paid           44484 Short Term       717    1152426      10+ years Home Mortgage
99995 Fully Paid           210584 Short Term       719    783389      1 year Home Mortgage
99996 Fully Paid           147070 Short Term       725    475437      7 years Own Home
99997 Fully Paid           99999999 short Term      732    1289416      1 year Rent Debt Consolidation
99998 Fully Paid           103136 Short Term       742    1150545      6 years Rent Debt Consolidation
99999 Fully Paid           530332 Short Term       746    1717524      9 years Rent Debt Consolidation
   Purpose Monthly.Debt
   small_business 6280.64
   other          3727.61
   other          2202.86
   Rent Debt Consolidation 13109.05
   Rent Debt Consolidation 7315.57
   Rent Debt Consolidation 9890.07
   Years.of.Credit.History Number.of.Open.Accounts Number.of.Credit.Problems Current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
99994           21.0             6            0         961932            0            0            0
99995           17.4             6            0          456         259160            0            0
99996           22.3             5            0         47766         658548            0            0
99997            9.4            22            0         153045         509234            0            0
99998           18.8            12            1         109554         537548            1            0
99999           15.0             8            0         404225         738254            0            0
> missing.data? <- colSums(is.na(credit.loans))> colSums(is.na(credit.loans)) > 0.7 %% sort(decreasing=T)
```

### Scaling:

The additional thing which we are going to do over here is Vector scaling. In many machine learning algorithms, to bring all features in the same standing, we need to do scaling so that one significant number doesn't impact the model just because of their large magnitude. scaling in machine learning is one of the most critical steps during the pre-processing of data before creating a machine learning model. Scaling can make a difference between a weak machine learning model and a better one. We are going to scale all are numeric columns.

## Data set after scaling:

Loan.Status	Current.Loan.Amount	Term	Credit.Score	Annual.Income	Years.in.current.job	Home.Ownership	Purpose	Monthly.Debt
Fully Paid	-0.3559827	Short Term	-0.2769929	-0.2167940	8 years	Home Mortgage	Home Improvements	-1.0889320
Fully Paid	-0.3617433	Short Term	0.0000000	0.0000000	10+ years	Home Mortgage	Debt Consolidation	1.2175321
Fully Paid	2.7763514	Short Term	-0.2528712	0.8779274	8 years	Own Home	Debt Consolidation	0.8811506
Fully Paid	-0.3590581	Long Term	-0.2679473	-0.5876091	3 years	Own Home	Debt Consolidation	-0.7992273
Fully Paid	-0.3644524	Short Term	0.0000000	0.0000000	5 years	Rent Debt	Debt Consolidation	0.1780034
Charged Off	-0.3634965	Short Term	4.6837898	-0.4951398	10+ years	Home Mortgage	Debt Consolidation	-0.1728758
Years.of.Credit.History	Number.of.Open.Accounts	Number.of.Credit.Problems	Current.Credit.Balance	Maximum.Open.Credit	Bankruptcies	Tax.Liens		
-0.1424297	-1.0237046	1	-0.176647582	-0.041035323	1	0		
0.4134949	4.7648980	0	-0.171899754	0.010731628	0	0		
-0.4702827	1.3715792	1	0.008921789	-0.001277969	0	0		
-0.8836626	-0.42448837	0	-0.101844033	-0.044588086	0	0		
-1.7246768	0.7727582	0	-0.109470864	-0.039791594	0	0		
-0.1281753	-1.0237046	0	-0.210892556	-0.058245493	0	0		

## Training and Testing the data:

We divide the data set using Hold out evaluation. We will divide the data set into 80% training data and 20% test data. The code can be seen below.

```
#test and train
credit.log=credit.log[sample(nrow(credit.log)),]
select.data=sample(1:nrow(credit.log),0.8*nrow(credit.log))
train.data.log=credit.log[select.data,]
test.data.log=credit.log[-select.data,]
head(train.data.log)

train.log.label=train.data.log$Loan.Status
test.log.label=test.data.log$Loan.Status
head(train.log.label)
head(test.log.label)

train.data.log=select(train.data.log,-1)
test.data.log=select(test.data.log,-1)
head(train.data.log)
head(test.data.log)

head(train.log.label)
head(test.log.label)
```

```

> head(test.data.log)
   Current.Loan.Amount Term Credit.Score Annual.Income Years.in.current.job Home.Ownership Purpose Monthly.Debt Years.of.Credit.History
59577      -0.3533046 Long Term   -0.3078989   0.20664578    10+ years Home Mortgage Debt Consolidation  0.98828392   2.0670143
39291       2.7763514 Short Term  -0.2528712   -0.74233683      3 years   Rent Debt Consolidation -0.53316522  -0.9264260
92098      -0.3638522 Short Term  -0.2724701   0.02797997      5 years   Rent Debt Consolidation  0.09891354  -0.8694081
57940      -0.3584697 short Term  0.0000000   0.00000000      3 years Home Mortgage Debt Consolidation  0.79216675   0.2424412
67977      -0.3598507 Short Term  -0.2611630   -0.47231558    10+ years Home Mortgage Debt Consolidation -0.27170687   2.1097777
62494      -0.3628631 short Term  -0.2739777   -0.78757486      4 years   Own Home Debt Consolidation -0.16309101  0.2139322
   Number.of.Open.Accounts Number.of.Credit.Problems Current.Credit.Balance Maximum.Open.Credit Bankruptcies Tax.Liens
59577           1.3715792          0            1.9840688   0.114341624      0      0
39291          -1.4229186          0            -0.7056264  -0.085528820      0      0
92098          -0.2252767          0            -0.1330585  -0.053210226      0      0
57940           1.1719722          0            -0.2436223  -0.006048222      0      0
67977           0.5731512          0            0.2636882  -0.016234207      0      0
62494          -0.2252767          0            -0.1622526  -0.036879063      0      0
> head(train.log.label)
[1] Fully Paid charged off Fully Paid Fully Paid Fully Paid charged off
Levels: Charged Off Fully Paid
> head(test.log.label)
[1] Charged Off Fully Paid Fully Paid Fully Paid Charged off Fully Paid
Levels: Charged Off Fully Paid

```

## Converting the dependent variable into binary

We converted our dependent variable into binary value so that we can calculate res value easily and we can classify the probability value close to 1 as “fully paid” and close to 0 as “Charged off”.

```

#converting dependent variable into binary for train data set
train.log.label=as.character(train.log.label)
head(train.log.label)
is.character(train.log.label)
str(train.log.label)
train.log.label[train.log.label == "Fully Paid"]=1
train.log.label[train.log.label == "charged off"]=0
train.log.label=as.integer(train.log.label)
head(train.log.label)

#converting dependent variable into binary for test data set

test.log.label=as.character(test.log.label)
head(test.log.label)
is.character(test.log.label)
str(test.log.label)
test.log.label[test.log.label == "Fully Paid"]=1
test.log.label[test.log.label == "charged off"]=0
test.log.label=as.integer(test.log.label)
head(test.log.label)

```

```

> train.log.label=as.character(train.log.label)
> head(train.log.label)
[1] "Fully Paid" "Charged off" "Fully Paid" "Fully Paid" "Fully Paid" "Charged off"
> is.character(train.log.label)
[1] TRUE
> str(train.log.label)
chr [1:79999] "Fully Paid" "Charged off" "Fully Paid" "Fully Paid" "Fully Paid" "Charged off" ...
> train.log.label[train.log.label == "Fully Paid"]=1
> train.log.label[train.log.label == "Charged off"]=0
> train.log.label=as.integer(train.log.label)
> head(train.log.label)
[1] 1 0 1 1 1 0
> test.log.label=as.character(test.log.label)
> head(test.log.label)
[1] "Charged off" "Fully Paid" "Fully Paid" "Charged off" "Fully Paid"
> is.character(test.log.label)
[1] TRUE
> str(test.log.label)
int [1:79999] 1 0 1 1 1 0 1 1 1 ...
> test.log.label[test.log.label == "Fully Paid"]=1
> test.log.label[test.log.label == "Charged off"]=0
> test.log.label=as.integer(test.log.label)
> head(test.log.label)
[1] 0 1 1 1 0 1

```

## Logistic Model

```

#logistic model
model.log.final=glm(train.log.label~.,data = train.data.log,family="binomial")
summary(model.log.final)

```

We calculated different cut off value and predicted accuracy of each model. Then we selected best Cut off value.

Res & Accuracy Calculation for 0.4

```

> for(i in 1:length(res)){
+   if(res[i]>0.4){
+     res[i]=1
+
+   }else{
+     res[i]=0
+   }
+ }
> library(Metrics)
Warning message:
package 'Metrics' was built under R version 3.6.3
> confmatrix=table(Actual_value=test.log.label,predicted_value= res)
> confmatrix
      predicted_value
Actual_value    0      1
      0  903  3753
      1      0 15344
> accuracy(test.log.label,res)
[1] 0.81235
```

```

Res & Accuracy Calculation for 0.2

```

> predict(model.log.final,test.data.log,type = "response")
> for(i in 1:length(res)){
+   if(res[i]>0.2){
+     res[i]=1
+   }
+   }else{
+     res[i]=0
+   }
+ }
> accuracy(test.log.label,res)
[1] 0.8192
>

```

### Res & Accuracy Calculation for 0.3

```

> res=predict(model.log.final,test.data.log,type = "response")
> for(i in 1:length(res)){
+   if(res[i]>0.3){
+     res[i]=1
+   }
+   }else{
+     res[i]=0
+   }
+ }
> accuracy(test.log.label,res)
[1] 0.8153
>

```

### Res & Accuracy Calculation for 0.5

```

glm.fit: fitted probabilities numerically 0 or 1 occurred
> # calculating res
> res=predict(model.log.final,test.data.log,type = "response")
> for(i in 1:length(res)){
+   if(res[i]>0.5){
+     res[i]=1
+   }
+   }else{
+     res[i]=0
+   }
+ }
> accuracy(test.log.label,res)
[1] 0.8222
>

```

### Res & Accuracy Calculation for 0.6

```
# calculating res
res=predict(model.log.final,test.data.log,type = "response")
for(i in 1:length(res)){
  if(res[i]>0.6){
    res[i]=1
  }else{
    res[i]=0
  }
}
accuracy(test.log.label,res)
[1] 0.8154
```

Res & Accuracy Calculation for 0.7

```
glm.fit: fitted probabilities numerically 0 or 1 o
> for(i in 1:length(res)){
+   if(res[i]>0.7){
+     res[i]=1
+
+   }else{
+     res[i]=0
+   }
+ }
> accuracy(test.log.label,res)
[1] 0.7358
> |
```

As we can see above the best cut off value for this model is **0.5**. This can be concluded by observing the accuracy for that cut off model. If we **increase** the cut off value, then the **accuracy decreases** and if we **decrease** the cut off value then also the **accuracy decreases**.

### Confusion Metrics for res 0.5:

```
> table(Actual_value=test.log.label,predicted_value= res)
           predicted_value
Actual_value      0      1
          0    930  3561
          1     19 15490
~ |
```

Precision: True positive/( True positive + False Positive)=930/(930+3561)

Precision = 0.2070=20.70%

Recall: True Positive/(True Positive + False Negative)= 930/(930+19)

Recall= 0.9799=97.99%

## C. Gradient Boosting Model

After performing data cleaning and scaling the data set we will divide the data set into training and testing as shown below. We will also convert our dependent variable into binary.

### Training and Testing

```
#test and train
credit.dummy1=credit.dummy1[sample(nrow(credit.dummy1)),]
select.data=sample(1:nrow(credit.dummy1),0.8*nrow(credit.dummy1))
train.data.grad=credit.dummy1[select.data,]
test.data.grad=credit.dummy1[-select.data,]
head(train.data.grad)

train.grad.label=train.data.grad$Loan.Status
test.grad.label=test.data.grad$Loan.Status
head(train.grad.label)
head(test.grad.label)

train.data.grad=select(train.data.grad,-1)
test.data.grad=select(test.data.grad,-1)
head(train.data.grad)
head(test.data.grad)

head(train.grad.label)
unique(test.grad.label)
```

### Gradient Boosting Model:

As we can see in the below screenshot, Credit score plays very important role whether the person will be able to pay the loan or not. Credit score has maximum rel.inference after that Current loan amount plays very important role and so on.

```

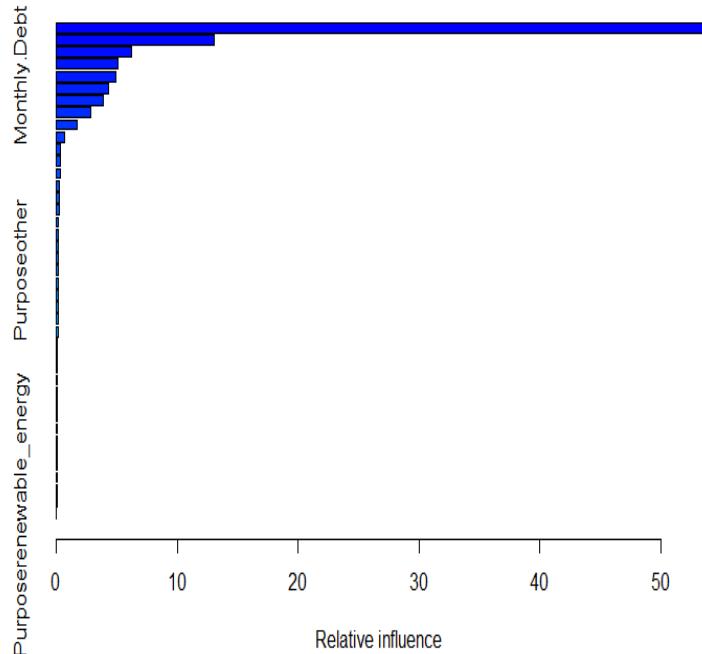
> Gradient.boost=gbm(train.grad.label~, data = train.data.grad,distribution = "gaussian",n.trees=100)
> summary(Gradient.boost)

var      rel.inf
Credit.Score          Credit.Score 53.99929468
Current.Loan.Amount   Current.Loan.Amount 13.01901695
Annual.Income          Annual.Income 6.17598785
Maximum.Open.Credit   Maximum.Open.Credit 5.08893058
Monthly.Debt          Monthly.Debt 4.87340883
Current.Credit.Balance Current.Credit.Balance 4.25571737
Years.of.Credit.History Years.of.Credit.History 3.87923690
`TermShort Term`     `TermShort Term` 2.84416431
Number.of.Open.Accounts Number.of.Open.Accounts 1.66851772
Home.Ownership.Rent   Home.Ownership.Rent 0.62553748
`Home.Ownership.Home Mortgage` `Home.Ownership.Home Mortgage` 0.34281448
Purposes.small.business Purposes.small.business 0.31983583
`PurposeDebt Consolidation` `PurposeDebt Consolidation` 0.29805864
Number.of.Credit.Problems Number.of.Credit.Problems 0.23161899
`Years.in.current.job10+ years` `Years.in.current.job10+ years` 0.22758145
Tax.Liens              Tax.Liens 0.20592560
Bankruptcies           Bankruptcies 0.17202490
`Years.in.current.job3 years` `Years.in.current.job3 years` 0.14591317
`Years.in.current.job< 1 year` `Years.in.current.job< 1 year` 0.13258425
`Years.in.current.job2 years` `Years.in.current.job2 years` 0.12768791
Purpose.other          Purpose.other 0.11448109
`Years.in.current.job7 years` `Years.in.current.job7 years` 0.11049969
`Years.in.current.job4 years` `Years.in.current.job4 years` 0.10736516
`Home.Ownership.own Home`   `Home.Ownership.own Home` 0.09974454
`PurposeBuy a Car`       `PurposeBuy a Car` 0.09528329
`Years.in.current.job5 years` `Years.in.current.job5 years` 0.09099810
`Years.in.current.job9 years` `Years.in.current.job9 years` 0.08425178
`Years.in.current.job1 year` `Years.in.current.job1 year` 0.08208809
Purpose.other          Purpose.other 0.07925212
`Years.in.current.job6 years` `Years.in.current.job6 years` 0.06871946
`PurposeHome Improvements` `PurposeHome Improvements` 0.06825031
`Years.in.current.job8 years` `Years.in.current.job8 years` 0.06387096
`PurposeMedical Bills`   `PurposeMedical Bills` 0.06375668
Purpose.major.purchase  Purpose.major.purchase 0.05153879
Purpose.evacation       Purpose.evacation 0.05016276
`PurposeBuy House`      `PurposeBuy House` 0.04983904
Purpose.moving          Purpose.moving 0.03437133
Purpose.wedding         Purpose.wedding 0.01886387
`PurposeTake a Trip`    `PurposeTake a Trip` 0.01823612
`PurposeEducational Expenses` `PurposeEducational Expenses` 0.01456893
Purpose.renewable.energy Purpose.renewable.energy 0.00000000

```

## Graph:

We have created a bar graph for the above model. The bar graph is based on the rel.inference value. As you can see Credit score has a maximum height compared to any other feature.



## Prediction Metrics:

```

n.trees = seq(from=100 ,to=10000, by=1) #no of trees-a vector of 100 values
#Generating a Prediction matrix for each Tree
predmatrix<-predict(Gradient.boost,test.data.grad,n.trees = n.trees)
dim(predmatrix) #dimentions of the Prediction Matrix
[1] 20000  9901

```

---

## Test Errors:

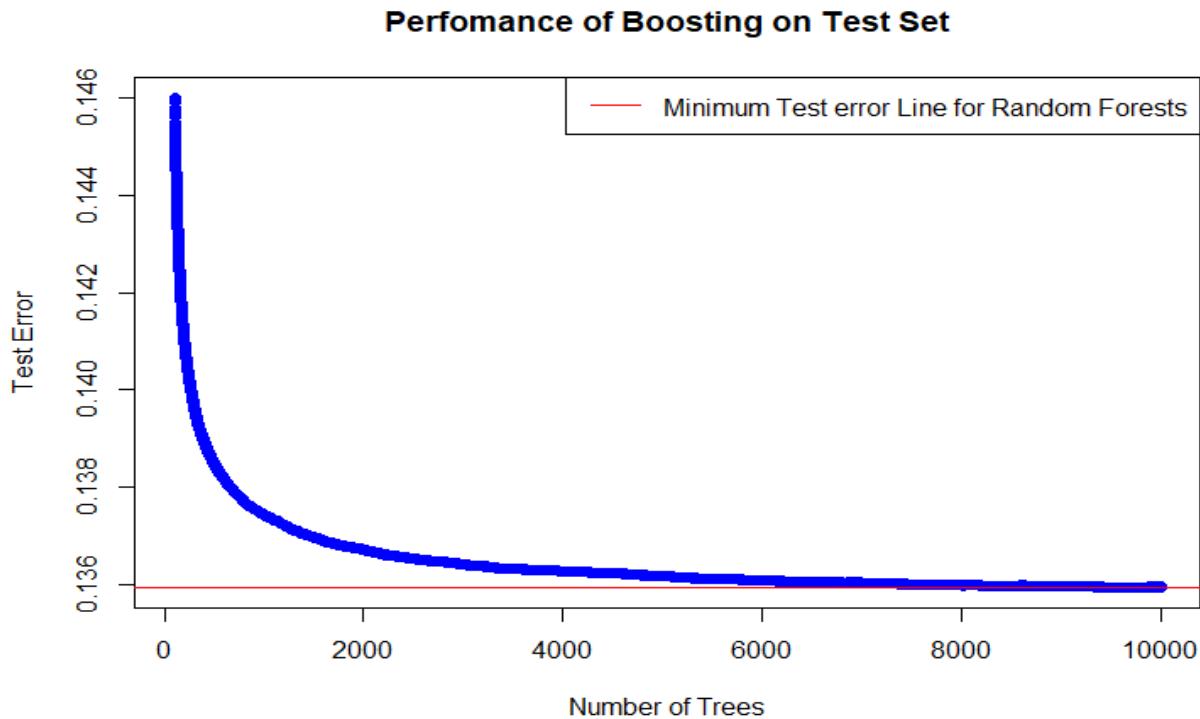
```

> test.error<-with(test.data.grad,apply( (predmatrix-test.grad.label)^2,2,mean))
> head(test.error) #contains the Mean squared test error for each of the 100 trees averaged
  100      101      102      103      104      105 
0.1460116 0.1459054 0.1458006 0.1456993 0.1456003 0.1455015
> 

```

## Graph:

We have created a graph for the test errors done by the gradient boosting model. As you can see in the graph as the number of trees increases the number of test errors decreases. At 10,000 trees you can see the test error is near to 0.135



## Accuracy

As the test error value for gradient boosting is 0.135. The accuracy of the model is 1- test error. So, the accuracy is 0.865. Therefore, the accuracy is 86.5%.

## 7. Conclusion:

- After comparing all the model, we found that Accuracy of KNN Model is 81.22%. Accuracy of Logistic Regression Model is 82.222%. Accuracy of Gradient Boosting Algorithm is 86.5%. So, the best model for data set is Gradient Boosting.
- We also calculated Precision and recall to find out the percentage of true value with the false positive and false negative.
- We used various data cleaning techniques for cleaning the data. We also understood the concept of scaling